

Chapter-08 Autoencoders

8.1 Introduction to Autoencoders:

Autoencoders are a type of unsupervised deep learning model that use hidden layers to decompose & then recreate their input. They have several applications:

- Dimensionality Reduction - Preprocessing for classification
- Identifying, 'essential' elements of input data, & filtering out noise.

Data Representation: Motivation:-

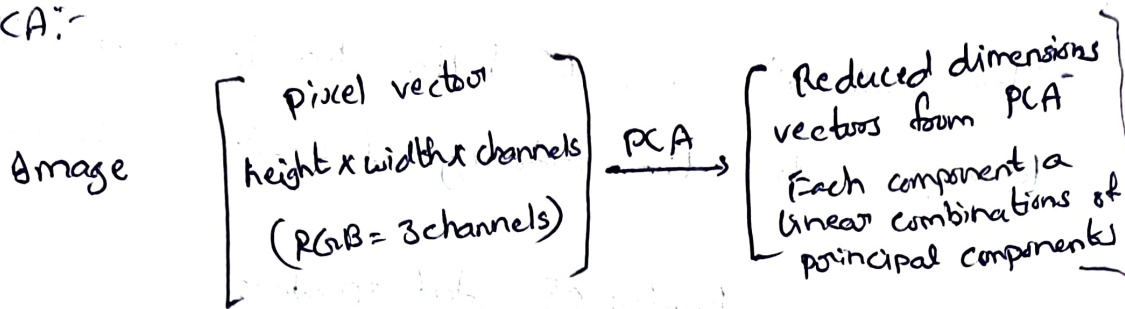
Suppose we want to find whether 2 images are 'similar':-
We could not look at every pixel wise & distinguish if 2 images are both or different.

But we want to ~~the~~ capture the image content of 2 images & check the similarity. Then we can say that 2 images are similar or different.

Non-Deep Learning Techniques:-

One approach could be Principal Component Analysis

PCA:-



PCA:- Find dimensions that capture most variance in original data


It usually reduces the dimensions of original image

- Two Dimensional eg:

- Direction of arrows represent principal components

- Length of arrows correspond to variance in original data

- Notice that both axes are combination of original features



Limitations of PCA:-

- Learned features are linear combinations of original features

- There may be complex, non-linear relationship between original features (pixels) & best lower dimensional representation

- The best representation can be defined in many different ways.

8.2 Autoencoders


Autoencoders are a neural networks architecture that focus the learning, of a lower dimensional representation of data, commonly images.

Autoencoders, step by step: Step 1: Feed image through encoder network.

Autoencoders, Step 2: Lower-dimensional embedding generated.

Autoencoders, Step 3: Embedding feed through decoder network

Autoencoders, Step 4: Reconstructed version of original data generated



Autoencoders, Step 5: Compare results to original (compute loss)

Results:

Network will learn lower dimensional space representing original data

Application: Autoencoders can be used to find image similarity.

The decoder network maps vectors from the lower-dimensional space to images

The decoder model can be used as a generative model once it is trained.

This is not oftenly done:-

- To get reasonable results, deep convolutional architecture is required
- Generally inferior to using GANs for image generation

Business Application:-

They have a wide variety of enterprise applications:-

- Dimensionality reduction as preprocessing for classification
- Information retrieval
- Anomaly detection
- Machine Translation
- Image related application (generation, denoising, preprocessing & compression)
- Drug discovery
- Popularity prediction for social media posts
- Sound & music synthesis

While most autoencoders use deep layers, autoencoders are often trained with a single layer each for the encoding & decoding step. An example is sparse autoencoders, which have been used successfully in recommender systems