

Chapter-09 Generative Adversarial Networks & Applications of Deep learning

Q.1 What is a Variational Autoencoder?

Variational Autoencoders also generate a latent representation & then use this representation to generate new samples (i.e. images)

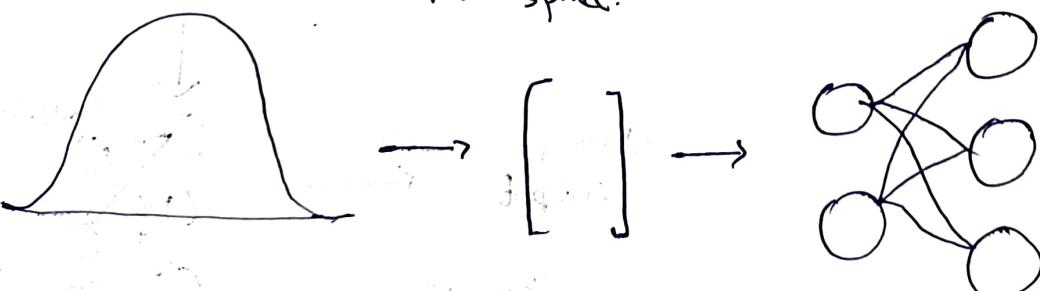
Important Features of variational autoencoders:-

- Data are assumed to be represented by a set of normally distributed latent factors
- The encoder generates parameters of these distributions namely μ & σ
- Images can be generated by sampling from these distributions

Goal of VAE :- generate image using the decoder

- latent vector: Each element drawn from a normal distribution
- Parameters learned by the decoder

Secondary goal:- have similar images be close together in latent space.



Working of VAE:

Variational autoencoders assume that the latent distributions is normally distributed, then learn to generate images from this distribution.

Step 1:- Image fed through encoder network

Step 2:- learning the parameters from normal dist
 (μ, σ)

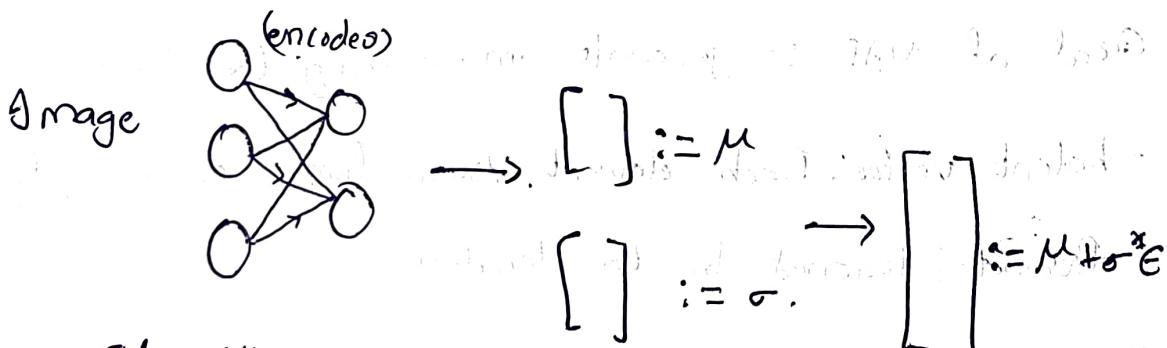
Step 3:- These are combined into one vector

- Random Noise $\epsilon \sim N(0,1)$ added.

combined vector $= \mu + \sigma * \epsilon$

Step 4:- This vector is fed through the decoder.

Step 5:- The deconstructed image is generated using the decoder network.



Step 01

Step 02.

Step 03.

Image
output

Step 05

↓
(decoder)
Step 04

q. Loss Function of VAE:

We want the VAE to learn to reconstruct the original image.

The VAE should do so from the space of vectors drawn from a standard normal distribution

The 2 components of the loss will be:

- A penalty for not reconstructing the image correctly
- A penalty for generating vectors of parameters $\mu \neq 0$ & $\sigma \neq 1$ (are different than 0 & 1 respectively, the parameters of standard normal distribution)

VAE have a loss function with 2 components:-

- 1) The pixel-wise difference between the reconstructed image & the original image. Many functions (like MSE, etc) can be used
- 2) The difference between the vectors produced by the encoder & the parameters of the standard normal distribution.
- * Specific loss used for this part of VAE is the "KL Divergence" between the generated data & the normal distribution.

→ $\mu = \begin{bmatrix} 0.7 \\ -0.6 \end{bmatrix} \longleftrightarrow \begin{bmatrix} 0 \\ 0 \end{bmatrix}$ → KL Divergence used to compare these vectors

→ vector interpreted as $\log(\sigma)$ since variance can't be negative $\left\{ \log(\sigma) = \begin{bmatrix} 1.0 \\ 0.6 \end{bmatrix} \longleftrightarrow \begin{bmatrix} 0 \\ 0 \end{bmatrix} \right\}$

$\log(\sigma)$ compared to 0 since $\log(0) = 0$

→ KL Divergence formula for

Normal distribution:-

$$KL = \frac{1}{2} \times \left(e^{\log(\sigma)} + (\log(\sigma) + 1) + (\mu_2^2) \right)$$

This function penalizes loss of $\log(\sigma)$ & μ respectively for being different from zero (0).

* Loss for $\log(\sigma)$: use the fact that $e^x + (x+1)$ is (strictly) convex positive definite function minimised at $x=0$

Note on KL Divergence:-

- It is not technically necessary to include it in VAE loss function to generate good latent space
- It helps generate a latent space where visually similar images are close in the latent space

VAE: Results

On the right:

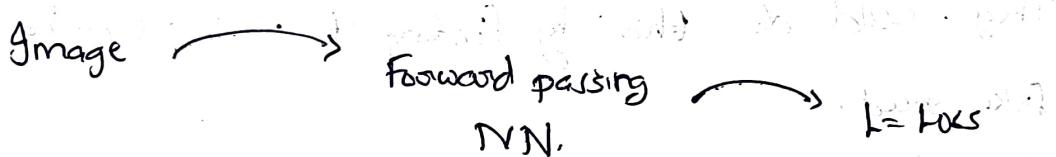
- VAE trained with 2 dimensions in the μ & σ vectors
- As the 2 dimensions of the μ vector vary from -15 to 15, we can plot the digits the trained decoder generates
- Close numbers in the latent space generates similar images

9.3 Generative Adversarial Networks (GAN)

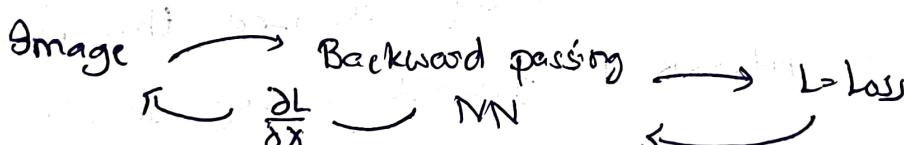
* Unsupervised learning & Powerful Generative Model.

Review: Adversarial Examples:

Neural network's differentiability makes them vulnerable to adversarial examples:



Adversarial examples: generated by repeatedly adding a function of $\frac{\partial L}{\partial x}$ to the input image



→ Take a training set & the focus on adjusting our original images in that backward pass in relation to each ~~one~~ one of our gradients.

GAN's Origin Story:

- The invention of GAN was connected to neural network's vulnerability to adversarial examples.
- Researchers were going to run a speech synthesis contest, to see which NN could generate the most realistic speech
 - A neural network - "the discriminator" - would judge whether the speech was real or not
 - They decided not to run the contest, bcz they realised people would generate speech to fool this particular network, rather than generating realistic speech
 - The researchers realized they could solve this by having the discriminator continually improve at distinguishing between real & fake speech
 - They could do this by feeding it real speech beside fake speech.
 - By feeding the gradient of resulting discriminator with respect to the input, $\frac{\partial L}{\partial x}$, back to a NN. generating the speech, they realized they could train a network to generate a very realistic speech.
 - As both networks try to beat the other, they both continuously improve.

GAN Definition:

GANs are a way of training a NN simultaneously

- One of the neural networks - "the generator" - learns to map random noise to "images" indistinguishable from those in some training set

Generator:



$$[] \rightarrow \text{conv1} \rightarrow \text{conv2} \rightarrow \text{conv3} \rightarrow \text{conv4} \rightarrow G(x)$$

Random
Noise

Discriminator: which are fake
which are from training set

$$G(x) \rightarrow \text{conv4} \rightarrow \text{conv3} \rightarrow \text{conv2} \rightarrow \text{conv1} \rightarrow D(x)$$

Image is indistinguishable from training set images

9.4 How GANs Work?

Training procedure of GAN:

Step 1: Randomly initialise weights of generator & discriminator networks

Step 2: Randomly initialise noise vector & generate image using generator

(generated img should be the same size of training set images as it is later fed through the Discriminator network)

Step 3:- Predict probability generated image is real using discriminating two.

(Output:- probability image is real)

Step 4:- Compute losses both assuming the image was fake : comparing $P_{\text{real}}^{x_G}$ to zero(0), & assuming it was real : comparing $P_{\text{real}}^{x_G}$ to 1(one):

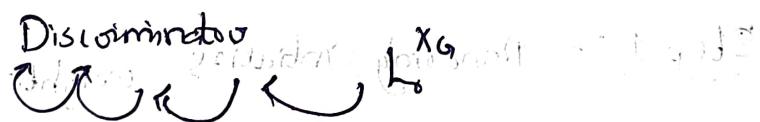
Generator $\rightarrow X_G \rightarrow$ Discriminator $\rightarrow P_{\text{real}}^{x_G}$

$$L_0^{x_G} = f(P_{\text{real}}^{x_G}, 0)$$

$$L_1^{x_G} = f(P_{\text{real}}^{x_G}, 1)$$

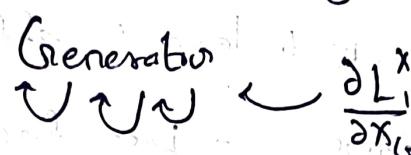
f = loss function.

Step 5:- Use $L_0^{x_G}$ to train the discriminator, we want to train the discriminator to output 0(O_0) for generated (fake) images.



Step 6:- Compute $\frac{\partial L_1^{x_G}}{\partial x_G}$. based on $L_1^{x_G}$, the penalty for the discriminator outputting a probability different than one(1), but ~~not~~ do not use this to train the discriminator.

Step 7 :- use $\frac{\partial L_1}{\partial x_G}$ to train the generator, we want it to generate images that the discriminator thinks are real. Here we continue to backpropagate through generator to update its weights.



Step 8 :- use the discriminator to calculate the probability that a real image is real.

Generator $\rightarrow x_R \rightarrow$ Discriminator $\rightarrow p_{\text{real}}^{x_R}$

Step 9 :- Compute $L_1^{x_R}$

Generator $\rightarrow x_R \rightarrow$ Discriminator $\rightarrow p_{\text{real}}^{x_R}$

$$L_1^{x_R} = f(\text{Preal}, 1)$$

Step 10 :- Use $L_1^{x_R}$ to train the discriminator by updating its weight.

→ Repeat this procedure with new random noise from generator each time. Continue until image from generator looks real. Values of losses from discriminator & generator may still be fluctuating when generator is producing realistic images, so can't use these alone to determine when to stop training.

* These are ways of quantifying image quality (e.g. Inception score)

9.5 Issues with Training GANs

- Training GANs effectively is highly dependent on both generator & discriminator learning at the same rate
- Ability of 2 networks to learn is affected by:
 - Network architecture
 - Loss functions
 - Learning rates
 - Optimization Techniques
- GANs are more sensitive than traditional neural networks to choice on these dimensions

What should you do to train a GAN?

compared with building a neural network for a supervised learning problem such as image classification or text generation, it is more important to read original papers & examine code on GitHub to see how researchers train their GANs.

- Famous examples of GANs include:
- Deepfakes → Replacing your face with any other celebrity
 - Age Interpolation → changing one's age
 - Text to Image
 - Lip sync

9.6 Additional Topics in Deep Learning

Hardware for AIs: GPUs

GPUs have become popular DL workhorse

- Feature thousands of small, simple cores specialized for numeric parallel computations
- Many transitions dedicated to computation
- GPUs excel at repeated similar instructions in parallel
- Optimized for parallel data throughout computation
- Major neural networks breakthrough since 2012 have been powered by GPU computations
- Performance has since increased $>5x$
- Once the data is in GPU memory the bottleneck are small
- NVIDIA® is highly popular in GPUs
- GPUs are great at parallelization.

CPU vs GPU:

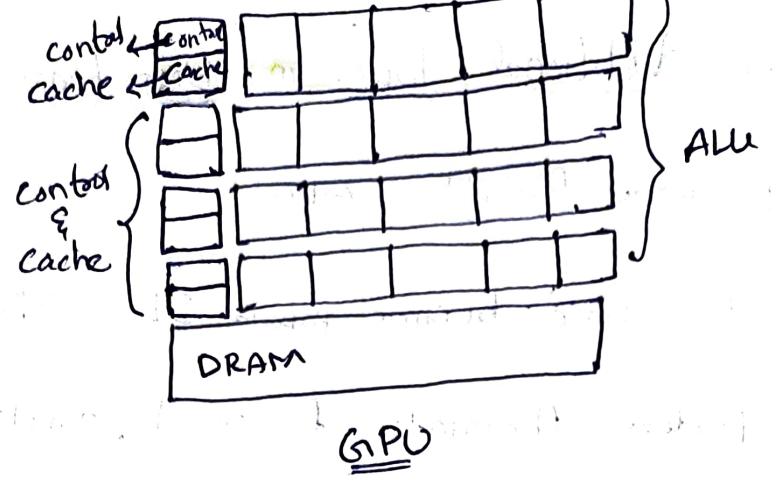
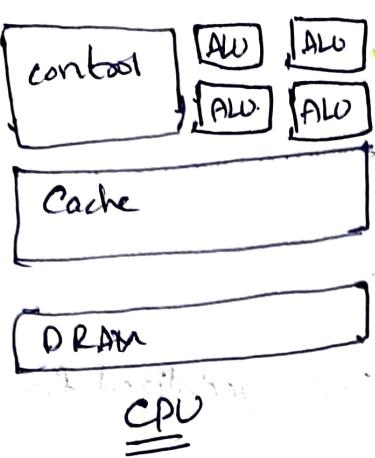
→ CPU:- Dozens of cores GPU: Thousands of cores

*ALU:- Arithmetic Logical Units → which performs arithmetic operations

*control:- Decode instructions into commands & calls on the ALU for any operations

*Cache:- serves as high speed memory, where instructions can be copied & delivered.

*DRAM:- for more, longer term memory.



- * GPU are specialised processors. More ALU & separate control & cache for ALUs
- * CPUs however are still main computation engines on most computers

Difference:

- CPU have dozen of cores (GPU have thousands) of less powerful cores
- CPU have fewer ALU & lower compute density than GPU
- CPU have lower latency & have larger cache memory, compared to GPU's
- GPU are designed for parallel tasks
- GPU perform well for a single instruction performed over a large amount of data
- GPU have additional overhead when copying data from main memory
- GPU can be better when a large number of memory swaps are needed.
- GPU are poor for task that can't be parallelized & for heavy processing on fewer data streams.
- CPU excels at serial tasks & are easy to program
- Popular programming lang, compilers, to be run on CPU by default.

Locally Interpretable Model-Agnostic Explanations:-

DL models are difficult to interpret:-

- Many parameters, Complex Network

One approach is to generate Locally Interpretable Model-Agnostic

Explanation (LIME):

- LIME treat the model as a black box & focus instead on the sensitivity of outputs to small changes in input
- Analogous to feature importance, LIME summarize the sensitivity of Regression or Classification outcomes to each variable
- Nonlinear & variable that cannot be perturbed present challenges to approach